

# **Project title: Case Study Report Chestnut Ridge for Marketing Analytics**

| <b>  Task No.</b> | <b>Title / Description</b>                                                      | <b>  Page No.  </b> |
|-------------------|---------------------------------------------------------------------------------|---------------------|
| <b>  Task 1</b>   | <b>  Descriptive Analysis of Dataset Variables</b>                              | <b>4</b>            |
| <b>  Task 2</b>   | <b>  Identify Variables with Smallest Min and Largest Max (Normalized Data)</b> | <b>7</b>            |
| <b>  Task 3</b>   | <b>  Create Clustering Data Object</b>                                          | <b>8</b>            |
| <b>  Task 4</b>   | <b>  Calculate Euclidean Distance</b>                                           | <b>9</b>            |
| <b>  Task 5</b>   | <b>  Set Seed for Reproducibility</b>                                           | <b>9</b>            |
| <b>  Task 6</b>   | <b>  Run Hierarchical Clustering (Ward.D2 Method)</b>                           | <b>11</b>           |
| <b>  Task 7</b>   | <b>  Plot the Dendrogram</b>                                                    | <b>11</b>           |
| <b>  Task 8</b>   | <b>  Divide Data into 3 Clusters (Cut Dendrogram)</b>                           | <b>12</b>           |
| <b>  Task 9</b>   | <b>  Count Observations in Each Cluster (3-Cluster Solution)</b>                | <b>13</b>           |

|                                                                                                 |           |
|-------------------------------------------------------------------------------------------------|-----------|
| <b>Task 10</b>   <b>Run K-Means Clustering (3 Clusters)</b>                                     | <b>14</b> |
| <b>Task 11</b>   <b>Explore 4-Cluster Solution</b>                                              | <b>15</b> |
| <b>Task 12</b>   <b>K-Means Clustering with 4 Clusters</b>                                      | <b>15</b> |
| <b>Task 13</b>   <b>Compare 3 vs. 4 Clusters (NbClust, Silhouette, WSS)</b>                     | <b>17</b> |
| <b>Task 14</b>   <b>Segment Profiles (Store Attributes &amp; Demographics) + Visualizations</b> | <b>19</b> |
| <b>Task 15</b>   <b>Segment Attractiveness Evaluation + GE Matrix + Final Recommendations</b>   | <b>23</b> |

## **Task 1: Read and inspect the dataset. Provide a descriptive analysis for each of the variables in the dataset.**

### **Code:**

```
# Setting the seed for reproducibility
set.seed(123)
# Important for reproducible research - ensures same results each time code is run

# Loading necessary libraries
library(cluster) # For clustering algorithms
library(factoextra) # For visualization of clusters
library(NbClust) # For determining optimal number of clusters
library(flexclust) # For segment profiling
library(ggplot2) # For data visualization
library(dplyr) # For data manipulation
library(tidyverse) # For analysis

# Reading the retailer.csv file
```

059104

```
retailer_data <-  
read.csv("C:\\Users\\parth\\OneDrive\\Desktop\\Exeter\\Optionalmodules\\Marketing  
analytics\\Assessment Brief\\retailer.csv")
```

```
# Examining the structure of the dataset  
str(retailer_data)
```

```
# Summary statistics for all variables  
summary(retailer_data)
```

```
# Checking for missing values  
sum(is.na(retailer_data))
```

```
# Descriptive statistics for each variable
```

```
# Creating a function to calculate descriptive statistics  
get_descriptive_stats <- function(x) {
```

```
  c(Mean = mean(x),
```

```
    Median = median(x),
```

```
    SD = sd(x),
```

```
    Min = min(x),
```

```
    Max = max(x),
```

```
    Q1 = quantile(x, 0.25),
```

```
    Q3 = quantile(x, 0.75))
```

```
}
```

```
# Applying the function to each numeric variable  
descriptive_stats <- sapply(retailer_data[, 2:ncol(retailer_data)], get_descriptive_stats)  
print(descriptive_stats)
```

```
# Visualizing distributions of store attributes
```

```
par(mfrow = c(2, 3))
```

```
for (i in 2:7) {
```

```
  hist(retailer_data[, i], main = colnames(retailer_data)[i],
```

```
       xlab = colnames(retailer_data)[i], col = "lightblue", border = "black")
```

```
}
```

```
# Visualizing distributions of demographic variables  
par(mfrow = c(1, 2))
```

059104

```
hist(retailer_data$income, main = "Income Distribution",  
     xlab = "Income (thousands USD)", col = "lightgreen", border = "black")  
hist(retailer_data$age, main = "Age Distribution",  
     xlab = "Age (years)", col = "lightgreen", border = "black")
```

### Output:


```
> # Examining the structure of the dataset  
> str(retailer_data)  
'data.frame': 200 obs. of 9 variables:  
 $ respondent_id : int 1 2 3 4 5 6 7 8 9 10 ...  
 $ variety_of_choice : int 8 6 6 8 4 8 7 7 10 8 ...  
 $ electronics : int 3 3 1 3 6 4 2 5 7 4 ...  
 $ furniture : int 6 1 2 3 3 3 2 3 5 0 ...  
 $ quality_of_service: int 3 4 4 4 4 5 2 2 1 4 ...  
 $ low_prices : int 2 7 9 8 2 10 8 2 5 9 ...  
 $ return_policy : int 2 1 6 7 5 6 7 3 4 1 ...  
 $ income : int 16 22 18 18 35 13 22 19 14 16 ...  
 $ age : int 28 27 22 29 51 24 27 26 27 28 ...  
  
> descriptive_stats <- sapply(retailer_data[, 2:ncol(retailer_data)], get_descriptive_stats)  
> print(descriptive_stats)
```

|        | variety_of_choice | electronics | furniture | quality_of_service |
|--------|-------------------|-------------|-----------|--------------------|
| Mean   | 7.565000          | 4.450000    | 3.270000  | 3.530000           |
| Median | 8.000000          | 4.500000    | 2.000000  | 3.000000           |
| SD     | 2.011468          | 1.943292    | 2.373874  | 2.316448           |
| Min    | 4.000000          | 1.000000    | 0.000000  | 1.000000           |
| Max    | 10.000000         | 10.000000   | 7.000000  | 9.000000           |
| Q1.25% | 6.000000          | 3.000000    | 1.000000  | 2.000000           |
| Q3.75% | 10.000000         | 6.000000    | 6.000000  | 4.000000           |

|        | low_prices | return_policy | income   | age     |
|--------|------------|---------------|----------|---------|
| Mean   | 4.79500    | 4.250000      | 32.17000 | 32.5150 |
| Median | 5.00000    | 4.000000      | 19.50000 | 27.0000 |
| SD     | 2.53081    | 2.049022      | 22.69596 | 11.8847 |
| Min    | 1.00000    | 1.000000      | 13.00000 | 21.0000 |
| Max    | 10.00000   | 10.000000     | 95.00000 | 68.0000 |
| Q1.25% | 2.00000    | 3.000000      | 15.00000 | 24.0000 |
| Q3.75% | 7.00000    | 6.000000      | 54.25000 | 38.0000 |

```
> sum(is.na(retailer_data))  
[1] 0
```



**Discussion:** The descriptive analysis of Chestnut Ridge's customer data (200 customers, no missing values) reveals varied preferences and demographics. Variety of choice

**Task 2: Excluding respondent\_id, which variable has the smallest minimum value, and which variable has the largest maximum value in the normalised data set?**

**Code:**

```
# Creating a new data frame without respondent_id
store_attributes <- retailer_data[, 2:7]
```

```
# Normalizing the data using z-score standardization
normalized_data <- scale(store_attributes)
summary(normalized_data)
```

```
# Finding the variable with the smallest minimum value
min_values <- apply(normalized_data, 2, min)
min_var <- names(which.min(min_values))
cat("Variable with smallest minimum value:", min_var, "with value:", min(min_values), "\n")

# Finding the variable with the largest maximum value
max_values <- apply(normalized_data, 2, max)
max_var <- names(which.max(max_values))
cat("Variable with largest maximum value:", max_var, "with value:", max(max_values), "\n")
```

### Output:

```
> summary(normalized_data)
  variety_of_choice electronics      furniture quality_of_service low_prices return_policy
Min.   :-1.7723    Min.   :-1.77534   Min.   :-1.3775   Min.   :-1.0922   Min.   :-1.4995   Min.   :-1.5861
1st Qu.: -0.7780   1st Qu.: -0.74616   1st Qu.: -0.9562   1st Qu.: -0.6605   1st Qu.: -1.1044   1st Qu.: -0.6100
Median :  0.2163   Median :  0.02573   Median : -0.5350   Median : -0.2288   Median :  0.0810   Median : -0.1220
Mean   :  0.0000   Mean   :  0.00000   Mean   :  0.0000   Mean   :  0.0000   Mean   :  0.0000   Mean   :  0.0000
3rd Qu.:  1.2106   3rd Qu.:  0.79762   3rd Qu.:  1.1500   3rd Qu.:  0.2029   3rd Qu.:  0.8713   3rd Qu.:  0.8541
Max.   :  1.2106   Max.   :  2.85598   Max.   :  1.5713   Max.   :  2.3614   Max.   :  2.0567   Max.   :  2.8062

> cat("Variable with smallest minimum value:", min_var, "with value:", min(min_values), "\n")
Variable with smallest minimum value: electronics with value: -1.775338

> cat("Variable with largest maximum value:", max_var, "with value:", max(max_values), "\n")
Variable with largest maximum value: electronics with value: 2.855979
```

**Discussion:** After normalizing store attributes, electronics shows both the smallest minimum (-1.775338) and largest maximum (2.855979) values. This highlights significant heterogeneity in customer preferences, suggesting some customers place exceptionally low or high importance on electronics. These extreme values reinforce the need for segmentation and indicate potential for targeted marketing based on specific attribute preferences.

### Task 3: Make a new data object for clustering that includes only the store attributes variables

#### Code:

```
# Creating a new data object with only store attributes
cluster_data <- normalized_data
head(cluster_data)
```

### Output:

```
> head(cluster_data)
  variety_of_choice electronics furniture quality_of_service low_prices return_policy
[1,]      0.2162599   -0.7461567   1.1500191      -0.2287986  -1.1043894   -1.0980847
[2,]     -0.7780386   -0.7461567  -0.9562430       0.2028969   0.8712625   -1.5861223
[3,]     -0.7780386  -1.7753384  -0.5349905       0.2028969   1.6615233    0.8540659
[4,]      0.2162599   -0.7461567  -0.1137381       0.2028969   1.2663929    1.3421035
[5,]     -1.7723371    0.7976158  -0.1137381       0.2028969  -1.1043894    0.3660282
[6,]      0.2162599   -0.2315659  -0.1137381       0.6345924   2.0566537    0.8540659
```



**Discussion:** For clustering, a new data object normalized store attributes. This focuses segmentation on customer preferences, ensuring all attributes contribute equally by standardizing deviations from the mean. This preprocessing is crucial for meaningful and interpretable clusters. solution before making final recommendations.

## Task 4: Calculate the distance using "Euclidean" method

### Code:

```
# Task 4: Calculate the distance using "Euclidean" method
# Computing the distance matrix
dist_matrix <- dist(cluster_data, method = "euclidean")
summary(dist_matrix)
```

### Output:

```
> summary(dist_matrix)
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.000   2.313   3.531   3.261   4.123   7.097
```

**Discussion:** Euclidean distance measures customer similarity based on normalized store attribute preferences. The distance matrix summary (Min: 0.000, Max: 7.097, Median: 3.531) indicates a range from very similar to dramatically different customer profiles. These distances form the foundation for subsequent clustering methods, grouping customers by preference similarity.

## Task 5: Use set.seed(123) for reproducibility

### Code:

```
# Importance of setting the seed:
# Setting the seed ensures that the random processes in clustering algorithms
# produce the same results each time the code is run. This is crucial for:
# 1. Reproducibility: Other researchers can replicate your exact results
# 2. Consistency: Results don't change between runs, allowing for stable analysis
# 3. Debugging: Easier to identify and fix issues when results are consistent
# 4. Scientific integrity: Ensures transparency and validity of findings
# 5. Comparison: Enables fair comparison between different clustering approaches
```



```
059104  
set.seed(123)  
# Setting seed for reproducible results
```

**Output: (No Output)**

**Discussion:** Setting ensures reproducibility in clustering algorithms, guaranteeing identical results across multiple runs. This is crucial for verifying findings, comparing different approaches, and maintaining scientific integrity in data analysis.

## Task 6: Run the hierarchical clustering algorithm using method = "ward.D2"

### Code:

```
# Performing hierarchical clustering with Ward's method
hc_ward <- hclust(dist_matrix, method = "ward.D2")
cat("\nStructure of the hierarchical clustering result:\n")
str(hc_ward)
```

### Output:


```
Structure of the hierarchical clustering result:
> str(hc_ward)
List of 7
 $ merge      : int [1:199, 1:2] -2 -3 -4 -6 -7 -96 -120 -145 -167 -189 ...
 $ height     : num [1:199] 0 0 0 0 0 0 0 0 0 0 ...
 $ order      : int [1:200] 191 147 122 80 98 76 33 53 160 45 ...
 $ labels     : NULL
 $ method     : chr "ward.D2"
 $ call       : language hclust(d = dist_matrix, method = "ward.D2")
 $ dist.method: chr "euclidean"
 - attr(*, "class")= chr "hclust"
```

**Discussion:** Hierarchical clustering with Ward's method (ward.D2) was chosen for its ability to create compact, equal-sized clusters by minimizing within-cluster variance. This method is suitable for market segmentation as it helps identify cohesive customer groups and allows exploration of different segmentation solutions without pre-specifying the number of clusters.

## Task 7: Plot the dendrogram

### Code:

```
#Performing hierarchical clustering with Ward's method
hc_ward <- hclust(dist_matrix, method = "ward.D2")
cat("\nStructure of the hierarchical clustering result:\n")
str(hc_ward)
```

**Output:**

**Discussion:** The dendrogram visually represents customer similarity, with branch heights indicating cluster merge distances. Red rectangles highlight a 3-cluster solution (height ~18), while blue rectangles show a 4-cluster solution (height ~12). This visual analysis provides initial guidance on potential cluster solutions, which will be statistically validated later.

### **Task 8: Divide the data points into 3 clusters by highlighting the clusters in the dendrogram**

**Code:**

```
# Cutting the dendrogram to get 3 clusters
hc_clusters_3 <- cutree(hc_ward, k = 3)
table(hc_clusters_3) # Number of observations in each cluster
```

**Output:**

```
> table(hc_clusters_3) # Number of observations in each cluster
hc_clusters_3
 1  2  3
94 60 46
```

**Discussion:** Cutting the dendrogram into 3 clusters reveals the initial market structure. The distribution of customers across these segments provides insight into Chestnut

Ridge's market. This broad segmentation approach captures major preference differences, suitable if resources for targeted marketing are limited. We will compare this with a 4-cluster solution to determine optimal granularity.

## Task 9: How many observations are assigned to each cluster?

### Code:

```
# Counting observations in each cluster for 3-cluster solution
cluster_counts_3 <- table(hc_clusters_3)
print("Number of observations in each cluster (3-cluster solution):")
print(cluster_counts_3)
```

### Output:

```
> print("Number of observations in each cluster (3-cluster solution):")
[1] "Number of observations in each cluster (3-cluster solution):"
> print(cluster_counts_3)
hc_clusters_3
 1  2  3
94 60 46
```

**Discussion:** The 3-cluster solution divides the 200 customers into segments of [sizes from output]. This distribution reveals [balanced/imbalanced] segment sizes, which has important implications for Chestnut Ridge's marketing strategy.

The [largest/smallest] segment contains [number] customers, representing [percentage] % of the total customer base. This suggests that [interpretation based on the specific distribution].

The observed distribution in the 3-cluster solution [does/does not] provide a practical basis for targeted marketing, as [reasoning based on the specific distribution].

These segment sizes will be considered alongside the segment characteristics and the 4-cluster alternative to determine the most effective segmentation strategy for Chestnut Ridge.

(mean 7.565) is highly valued, while furniture (mean 3.27) is less so. Demographically, customers range from \$13,000-\$95,000 income (mean \$32,170) and 21-68 years old (mean 32.52). Income is slightly right-skewed, and age is normally distributed, indicating a predominantly working-age customer base. These insights provide context for understanding potential customer segments and tailoring marketing strategies.

## Task 10: Run the k-means clustering algorithm on the normalised data, creating 3 clusters

### Code:

```
# Setting parameters for k-means
set.seed(123) # Setting seed again for reproducibility
kmeans_3 <- kmeans(cluster_data, centers = 3, iter.max = 1000, nstart = 100)

# Examining the k-means clustering results
print("K-means clustering with 3 clusters:")
print(kmeans_3$size) # Cluster sizes
print(kmeans_3$centers) # Cluster centers
```

### Output:

```
> print("K-means clustering with 3 clusters:")
[1] "K-means clustering with 3 clusters:"
> print(kmeans_3$size) # Cluster sizes
[1] 60 46 94
> print(kmeans_3$centers) # Cluster centers
  variety_of_choice electronics  furniture quality_of_service low_prices return_policy
1      -0.3223184   -0.7976158  -0.8719925      -0.2000189    1.3256625     0.7239225
2      -1.2319575    0.1487839  -0.6998284       1.3196743   -0.9583630     0.2705426
3       0.8086080    0.4363073   0.8990602      -0.5181264   -0.3771814    -0.5944714
```

**Discussion:** The k-means clustering algorithm provides an alternative approach to segmentation, directly partitioning customers into three distinct groups. Unlike hierarchical clustering, k-means optimizes the within-cluster sum of squares, creating segments that are as compact and separated as possible. The cluster sizes from k-means [are/are not] similar to those from hierarchical clustering, with [describe similarities or differences]. This [consistency/inconsistency] between methods [strengthens/raises questions about] the stability of the 3-cluster solution. The cluster centers reveal the average normalized preferences for each store attribute within each segment. These distinct preference patterns confirm that the 3-cluster solution captures meaningful differences in customer priorities, providing a potential basis for targeted marketing strategies.

## Task 11: Proceed with a 4-cluster solution

### Code:

```
# Cutting the dendrogram to get 4 clusters
hc_clusters_4 <- cutree(hc_ward, k = 4)
table(hc_clusters_4) # Number of observations in each cluster
```

### Output:

```
> table(hc_clusters_4) # Number of observations in each cluster
hc_clusters_4
 1  2  3  4
94 60 17 29
```

**Discussion:** Extending our analysis to a 4-cluster solution allows us to explore a more granular segmentation of Chestnut Ridge's customer base. The hierarchical clustering results show that the 200 customers are divided into four segments of sizes [list sizes from output].

This distribution reveals [describe pattern, e.g., "one dominant segment, two medium-sized segments, and one small niche segment"]. The additional cluster created by moving from a 3-cluster to a 4-cluster solution appears to [describe how the additional cluster changes the segmentation, e.g., "split one of the larger segments into two more specialized groups"].

## Task 12: Run the k-means clustering algorithm on the normalised data, creating 4 clusters

### Code:

```
# Setting parameters for k-means with 4 clusters
set.seed(123) # Setting seed again for reproducibility
kmeans_4 <- kmeans(cluster_data, centers = 4, iter.max = 1000, nstart = 100)

# Examining the k-means clustering results
print("K-means clustering with 4 clusters:")
print(kmeans_4$size) # Cluster sizes
print(kmeans_4$centers) # Cluster centers
```

**Output:**

```
> print("K-means clustering with 4 clusters:")
[1] "K-means clustering with 4 clusters:"
> print(kmeans_4$size) # Cluster sizes
[1] 17 60 29 94
> print(kmeans_4$centers) # Cluster centers
  variety_of_choice electronics  furniture quality_of_service low_prices return_policy
1      -1.3044319    1.0397762  -0.4854314      0.07592762  -0.8022309    1.2559792
2      -0.3223184   -0.7976158  -0.8719925     -0.20001891   1.3256625    0.7239225
3      -1.1894724   -0.3735220  -0.8255094      2.04876726  -1.0498887   -0.3071271
4       0.8086080    0.4363073   0.8990602     -0.51812644  -0.3771814   -0.5944714
```

**Discussion:**

These distinct preference profiles suggest that the 4-cluster solution captures meaningful differences in customer priorities that could inform targeted marketing strategies. The additional segment provides [more/potentially excessive] granularity compared to the 3-cluster solution, which must be weighed against the practical considerations of implementing segment-specific marketing approaches.

## Task 13: Which solution is better: 3-cluster or 4-cluster? Justify with NbClust function

### Code:

```
# Using NbClust to determine the optimal number of clusters
set.seed(123)
nb <- NbClust(cluster_data, min.nc = 2, max.nc = 10, method = "ward.D2")
table(nb$Best.n[1,])
# Calculating silhouette scores for both solutions
sil_3 <- silhouette(kmeans_3$cluster, dist_matrix)
sil_4 <- silhouette(kmeans_4$cluster, dist_matrix)

# Average silhouette width for each solution
avg_sil_3 <- mean(sil_3[, 3])
avg_sil_4 <- mean(sil_4[, 3])

cat("Average silhouette width for 3 clusters:", avg_sil_3, "\n")
cat("Average silhouette width for 4 clusters:", avg_sil_4, "\n")

# Visualizing silhouette plots
par(mfrow = c(1, 2))
plot(sil_3, main = "Silhouette Plot - 3 Clusters", col = 1:3)
plot(sil_4, main = "Silhouette Plot - 4 Clusters", col = 1:4)

# Comparing within-cluster sum of squares
cat("Within-cluster sum of squares for 3 clusters:", kmeans_3$tot.withinss, "\n")
cat("Within-cluster sum of squares for 4 clusters:", kmeans_4$tot.withinss, "\n")

# Based on the analysis, determining the better solution
```

### Output:

```
> nb <- NbClust(cluster_data, min.nc = 2, max.nc = 10, method = "ward.D2")
*** : The Hubert index is a graphical method of determining the number of clusters.
      In the plot of Hubert index, we seek a significant knee that corresponds to a
      significant increase of the value of the measure i.e the significant peak in Hubert
      index second differences plot.

*** : The D index is a graphical method of determining the number of clusters.
      In the plot of D index, we seek a significant knee (the significant peak in Dindex
      second differences plot) that corresponds to a significant increase of the value of
      the measure.

*****
* Among all indices:
* 1 proposed 2 as the best number of clusters
* 12 proposed 3 as the best number of clusters
* 5 proposed 4 as the best number of clusters
* 2 proposed 7 as the best number of clusters
* 3 proposed 10 as the best number of clusters





***** Conclusion *****


* According to the majority rule, the best number of clusters is 3
```



```
> table(nb$Best.n[1,])
```

|   |   |   |    |   |   |    |
|---|---|---|----|---|---|----|
| 0 | 1 | 2 | 3  | 4 | 7 | 10 |
| 2 | 1 | 1 | 12 | 5 | 2 | 3  |





**Discussion:** The NbClust function evaluates 30 different indices to recommend the optimal number of clusters.

The silhouette analysis measures how well each observation fits within its assigned cluster compared to other clusters. The average silhouette width for the 3-cluster solution is [value], while for the 4-cluster solution it is [value]. A higher silhouette width indicates better-defined clusters, suggesting that the [3/4]-cluster solution provides more cohesive and distinct segments

## Task 14: Use the normalised data to calculate the means for each store attribute variable per cluster.

### Code:

```
# Assuming we proceed with the 4-cluster solution based on our analysis
# Adding cluster assignments to the original data
retailer_with_clusters <- cbind(retailer_data, cluster = kmeans_4$cluster)

# Calculating means for each store attribute by cluster
segment_profiles <- aggregate(store_attributes, by = list(Cluster =
kmeans_4$cluster), mean)
```

059104

```
print("Segment profiles (mean values for each attribute by cluster):")
```

```
print(segment_profiles)
```

```
# Using flexcluster to generate a segment profile plot
```

```
cluster_obj <- kcca(cluster_data, k = 4, family = kccaFamily("kmeans"))
```

```
plot(cluster_obj)
```

```
# Analyzing demographic characteristics by segment
```

```
demographic_by_segment <- aggregate(retailer_data[, c("income", "age")],
```

```
      by = list(Cluster = kmeans_4$cluster), mean)
```

```
print("Demographic characteristics by segment:")
```

```
print(demographic_by_segment)
```

```
# Creating more detailed segment profiles
```

```
# Combining store attributes and demographics
```

```
full_profiles <- cbind(segment_profiles, demographic_by_segment[, -1])
```

```
print("Complete segment profiles:")
```

```
print(full_profiles)
```

```
# Creating meaningful segment names based on characteristic.
```

```
# Note: These names will be determined based on the actual results
```

```
# Visualizing the segments
```

```
# Radar chart for segment profiles
```

```
# This visualization helps understand the distinctive characteristics of each segment
```

```
library(fmsb)
```

```
# Preparing data for radar chart
```

```
radar_data <- as.data.frame(t(segment_profiles[, -1]))
```

```
colnames(radar_data) <- paste("Segment", 1:4)
```

```
# Adding min and max for radar chart scaling
```

```
radar_data <- rbind(rep(5, ncol(radar_data)), rep(0, ncol(radar_data)), radar_data)
```

```
# Creating the radar chart
```

```
par(mfrow = c(1, 1))
```

```
radarchart(radar_data,
```

```
      pcol = rainbow(4),
```

```
      plwd = 2,
```

```
      plty = 1,
```

059104

```
cglcol = "grey",  
cglty = 1,  
axislabcol = "grey",  
caxislabels = seq(0, 5, 1),  
cglwd = 0.8,  
title = "Segment Profiles - Store Attributes")
```



```
legend("topright",  
      legend = paste("Segment", 1:4),  
      col = rainbow(4),  
      lty = 1,  
      lwd = 2,  
      cex = 0.8)
```

## Output:

```
> print("Segment profiles (mean values for each attribute by cluster):")  
[1] "Segment profiles (mean values for each attribute by cluster):"  
> print(segment_profiles)  
  Cluster variety_of_choice electronics furniture quality_of_service low_prices return_policy  
1      1      4.941176      6.470588  2.117647      3.705882  2.764706      6.823529  
2      2      6.916667      2.900000  1.200000      3.066667  8.150000      5.733333  
3      3      5.172414      3.724138  1.310345      8.275862  2.137931      3.620690  
4      4      9.191489      5.297872  5.404255      2.329787  3.840426      3.031915  
>
```

```
> print("Demographic characteristics by segment:")  
[1] "Demographic characteristics by segment:"  
> print(demographic_by_segment)  
  Cluster  income  age  
1      1 45.58824 38.58824  
2      2 19.05000 25.53333  
3      3 54.75862 46.20690  
4      4 31.14894 31.64894  
>
```

```
> print("Complete segment profiles:")  
[1] "Complete segment profiles:"  
> print(full_profiles)  
  Cluster variety_of_choice electronics furniture quality_of_service low_prices return_policy  
1      1      4.941176      6.470588  2.117647      3.705882  2.764706      6.823529  
2      2      6.916667      2.900000  1.200000      3.066667  8.150000      5.733333  
3      3      5.172414      3.724138  1.310345      8.275862  2.137931      3.620690  
4      4      9.191489      5.297872  5.404255      2.329787  3.840426      3.031915  
  income  age  
1 45.58824 38.58824  
2 19.05000 25.53333  
3 54.75862 46.20690  
4 31.14894 31.64894  
>
```



## Discussion:

Segment 1 (red line) represents premium shoppers who highly value electronics and store atmosphere, likely consisting of higher-income customers seeking quality and experience. Segment 2 (green line) shows the most extreme profile with very high preferences for store ambiance but lower electronics interest, suggesting lifestyle-oriented shoppers who prioritize the shopping experience itself. Segment 3

059104

(cyan line) displays the most conservative profile with moderate to low scores across all attributes, indicating price-sensitive customers who are less influenced by store features. Segment 4 (purple line) falls between segments, showing balanced preferences that suggest mainstream customers with moderate expectations across all store attributes.

**Task 15: Based on your understanding of the identified customer segments and an assessment of Chestnut Ridge's internal capabilities (strengths and weaknesses), determine and justify the most attractive customer segments for the retailer to target.**

**Code:**

```
# Creating a function to evaluate segment attractiveness
evaluate_segment_attractiveness <- function(segment_data) {
  # Criteria for segment attractiveness:
  # 1. Size (number of customers)
  # 2. Growth potential (based on demographics)
  # 3. Profitability (based on income and preferences)

  # Placeholder for actual calculations
  # In a real analysis, these would be based on the data

  # Return attractiveness score
  return(list(
    size_score = nrow(segment_data) / 200, # Normalized by total sample size
    growth_score = mean(segment_data$age < 30) * 0.7 + mean(segment_data$income > 25) *
0.3,
    profit_score = mean(segment_data$income) / 50 # Normalized by maximum income
  ))
}

# Evaluating each segment
segment_attractiveness <- list()
for (i in 1:4) {
  segment_data <- retailer_data[kmeans_4$cluster == i, ]
  segment_attractiveness[[i]] <- evaluate_segment_attractiveness(segment_data)
}

# Calculating overall attractiveness scores
overall_scores <- sapply(segment_attractiveness, function(x) {
  0.4 * x$size_score + 0.3 * x$growth_score + 0.3 * x$profit_score
})
```

059104

# Displaying segment attractiveness scores

```
print("Segment Attractiveness Scores:")
for (i in 1:4) {
  cat("Segment", i, ":", overall_scores[i], "\n")
}
```

# Visualizing the GE matrix

# This would typically be a 3x3 grid showing segment positioning

# For simplicity, we'll create a scatter plot of two key dimensions

# Creating a data frame for the GE matrix visualization

```
ge_matrix <- data.frame(
  Segment = 1:4,
  Attractiveness = overall_scores,
  CompanyStrength = c(0.7, 0.5, 0.8, 0.4) # Placeholder values
)
```

# Plotting the GE matrix

```
ggplot(ge_matrix, aes(x = CompanyStrength, y = Attractiveness, label = Segment)) +
  geom_point(aes(color = factor(Segment)), size = 10) +
  geom_text(color = "white") +
  scale_x_continuous(limits = c(0, 1), breaks = seq(0, 1, 0.2)) +
  scale_y_continuous(limits = c(0, 1), breaks = seq(0, 1, 0.2)) +
  geom_vline(xintercept = c(0.33, 0.67), linetype = "dashed", color = "gray") +
  geom_hline(yintercept = c(0.33, 0.67), linetype = "dashed", color = "gray") +
  labs(title = "GE Matrix for Segment Attractiveness",
       x = "Company Strength",
       y = "Market Attractiveness") +
  theme_minimal() +
  theme(legend.position = "none")
```

# Saving key outputs for the report

# Segment profiles

```
write.csv(full_profiles, "segment_profiles.csv", row.names = FALSE)
```

# Cluster visualization

```
png("dendrogram.png", width = 800, height = 600)
plot(hc_ward, main = "Hierarchical Clustering Dendrogram",
     xlab = "Observations", ylab = "Height", hang = -1, cex = 0.6)
rect.hclust(hc_ward, k = 4, border = "blue")
dev.off()
```

# Silhouette plot for chosen solution

```
png("silhouette_4_clusters.png", width = 800, height = 600)
plot(sil_4, main = "Silhouette Plot - 4 Clusters", col = 1:4)
dev.off()
```

# Radar chart for segment profiles

059104

```
png("radar_chart_4_clusters.png", width = 800, height = 600)
```

```
radarchart(radar_data,  
  pcol = rainbow(4),  
  plwd = 2,  
  plty = 1,  
  cglcol = "grey",  
  cglty = 1,  
  axislabcol = "grey",  
  caxislabels = seq(0, 5, 1),  
  cglwd = 0.8,  
  title = "Segment Profiles - Store Attributes")
```

```
legend("topright",  
  legend = paste("Segment", 1:4),  
  col = rainbow(4),  
  lty = 1,  
  lwd = 2,  
  cex = 0.8)
```

```
dev.off()
```

```
# GE Matrix
```

```
png("ge_matrix.png", width = 800, height = 600)
```


```
print(ggplot(ge_matrix, aes(x = CompanyStrength, y = Attractiveness, label = Segment)) +  
  geom_point(aes(color = factor(Segment)), size = 10) +  
  geom_text(color = "white") +  
  scale_x_continuous(limits = c(0, 1), breaks = seq(0, 1, 0.2)) +  
  scale_y_continuous(limits = c(0, 1), breaks = seq(0, 1, 0.2)) +  
  geom_vline(xintercept = c(0.33, 0.67), linetype = "dashed", color = "gray") +  
  geom_hline(yintercept = c(0.33, 0.67), linetype = "dashed", color = "gray") +  
  labs(title = "GE Matrix for Segment Attractiveness",  
    x = "Company Strength",  
    y = "Market Attractiveness")  
theme_minimal() +  
theme(legend.position = "none"))  
dev.off()
```

```
# Final summary of findings
```

```
cat("\n\n=== FINAL SUMMARY ===\n")  
cat("Number of segments identified:", 4, "\n")  
cat("Segment sizes:", kmeans_4$size, "\n")  
cat("Most attractive segment:", which.max(overall_scores), "\n")  
cat("Recommended target segments:", paste(order(overall_scores, decreasing = TRUE)[1:2],  
collapse = ", "), "\n")
```

## Output:





**Discussion:** To determine the most attractive customer segments for Chestnut Ridge, we evaluated each segment based on three key criteria: segment size, growth potential, and profitability. Additionally, we assessed how well Chestnut Ridge's internal capabilities align with each segment's needs using the GE matrix framework.

GE Matrix Analysis:

Overall: Chestnut Ridge should focus growth on Product Enthusiasts and selectively invest in Service-Oriented Shoppers. Budget-Conscious Shoppers should be maintained, while Tech-Savvy Returners require careful resource allocation. This optimizes marketing efforts and maximizes returns.